Question 1: Linear Regression — Predicting House Prices

Dataset: House Prices - Advanced Regression Techniques | Kaggle

Objective: Build a Linear Regression model from scratch and compare its performance with a prebuilt model from sklearn.

1. **Preprocess the Data**
    - Download and load train.csv.
    - Use the feature OverallQual to predict the target SalePrice.
    - Split the dataset into training and test sets (80% train, 20% test) to evaluate model performance using MSE.
2. **Train Linear Regression from Scratch**
    - Implement a function that calculates the slope and intercept using the least squares formula.
    - Train this model using the train set.
3. **Test Your Custom Model**
    - Use the same feature (OverallQual) from the test set to make predictions.
    - Calculate the Mean Squared Error (MSE) on the test predictions.
4. **Train and Test Sklearn's Linear Regression Model**
    - Use sklearn.linear_model.LinearRegression to train a model on the train set using the same feature.
    - Predict on the test set and compute the MSE for this model as well.
5. **Create the Following Plots**
    - Plot 1 (Train Comparison): Plot the custom model vs sklearn model predictions on train set.
    - Plot 2 (Test Comparison): Plot the custom model vs sklearn model predictions on test set.
    - In each plot, include both, the data points and the fitted lines. Add clear titles, axis labels and legends.

Note: Use only train.csv for training/testing split; ignore Kaggle's test.csv in this assignment.

Question 2: Logistic Regression — Predicting Student Pass/Fail Outcome

Dataset: Students Performance in Exams

Objective: Build a Logistic Regression model from scratch and compare its performance with a prebuilt model from sklearn.

1. **Preprocess the Data**
    - Download and load StudentsPerformance.csv (from the Kaggle dataset).
    - Create a new column PassedMath (1 if math score ≥ 50 and 0 otherwise)
    - Use "reading score" as the only feature to predict, PassedMath.
    - Split the dataset into training and test sets (80% train, 20% test).
2. **Train Logistic Regression from Scratch**
    - Implement logistic regression using the sigmoid function and gradient descent (for 1000 iterations, with a learning rate of 0.01).
    - Train the model on the training data.
3. **Test Your Custom Model**
    - Use your model to predict outcomes on the test set.
    - If a predicted probability is ≥ 0.5, classify it as 1; otherwise, classify as 0.
    - Calculate accuracy of your model and create a confusion matrix (you may use sklearn.metrics.confusion_matrix for visualization).
4. **Train and Test Sklearn's LogisticRegression Model**
    - Use sklearn.linear_model.LogisticRegression to train a model using the same feature.
    - Predict on the test set using the sklearn model, calculate accuracy, and plot the confusion matrix.

5. Additional Model Evaluation

- Calculate and compare accuracy, precision, recall, and F1-score for both your custom logistic regression model and sklearn's model on the test set.