

Neural Network with PyTorch — Classifying Iris Flowers

Dataset: Iris Dataset (download the iris.csv)

Objective: Build and train a neural network using PyTorch to classify iris flowers into one of three species.

1. Load & Preprocess the Data

- Load the dataset from the iris.csv file into a Pandas DataFrame.
- Encode the target column (species) as integer labels (0, 1, 2).
- Split the dataset into 80% training and 20% testing.
- Normalize the feature values (e.g., using StandardScaler).

2. Visualize the Data

- Create box plots for each feature (sepal length, sepal width, petal length, petal width), separated by species to show the distribution and spread per class.
- Optional: Create a pairplot (scatterplot matrix) of all features, colored by species.

3. Build the Neural Network

- Use PyTorch to define the following architecture:
 - Input layer: 4 neurons (one for each feature)
 - Hidden layer: 16 neurons with ReLU activation
 - Output layer: 3 neurons
- Use CrossEntropyLoss as the loss function.
- Use an optimizer such as SGD or Adam.

4. Train the Model

- Train the model for 100 epochs.
- After every 10 epochs:
 - Calculate and print training and test loss.
 - Calculate and print training and test accuracy.

5. Create the Following Plots

- Plot 1: Training and test loss vs. epochs
- Plot 2: Training and test accuracy vs. epochs
- Plot 3: Final confusion matrix (visualized as a heatmap with class labels)

Optional Extension: Understanding and Implementing Backpropagation from Scratch

Objective: Gain a deeper understanding of how neural networks learn by manually implementing the backpropagation algorithm. This will help you understand how neural networks optimize weights and biases using gradients.

1. Understand the Theory of Backpropagation

- Read about how the backpropagation algorithm works:
 - It computes the gradient of the loss function with respect to each weight using the chain rule.
 - These gradients are then used to update the weights in a direction that reduces the loss (typically by moving opposite to the gradient).

2. Implement Backpropagation from Scratch

- Instead of using PyTorch's built-in autograd and optimizer, implement backpropagation manually to compute gradients and update weights using gradient descent.
- Manually calculate:
 - Forward Pass: Calculate the activations of each layer.
 - Backward Pass: Compute the gradient of the loss with respect to each weight and bias, using the chain rule.
 - Weight Update: Manually update the weights using the computed gradients and a learning rate.

3. Compare with PyTorch's Built-in Backpropagation

- Once you've implemented backpropagation, compare the results (accuracy, loss) with the PyTorch model that uses the built-in autograd system and an optimizer.
- Verify that both implementations are consistent and produce similar results.

4. Experiment with Different Learning Rates

- After implementing backpropagation manually, experiment with different learning rates to observe how the network's training dynamics change.